Debugging LLM Facilitation through Synthetic Discussion Generation

No Author Given

No Institute Given

Abstract. Human facilitation is insufficiently scalable for modern online spaces. While Large Language Models (LLMs) may be able to automate the facilitation of online discussions, no large-scale studies have tested this hypothesis due to significant costs associated with human involvement. An effective solution is synthetic discussion simulations to create initial pilot experiments. We establish design principles and create a new Synthetic Discussion Generation (SDG) framework, which we use to prototype the development of LLM facilitators by generating data without human involvement. Through fast and cheap experiments, we find severe limitations on the capabilities of LLM facilitators; they cannot utilize effectively human facilitation strategies, nor decide when to intervene in a discussion. We develop an open-source python framework and release a large, publicly available dataset containing LLM-generated discussions using multiple open-source models. These resources can be used for LLM facilitator finetuning as well as behavioral analysis of current out-of-the-box LLMs in the task.

1 Introduction

The modern social media environment has evolved to be extremely demanding, with users facing ever-increasing threats such as targeted misinformation [15,17], hate speech [25], and polarization [41]. These threats can cause serious emotional and mental harm [45], radicalization [14], real-world violence [44], sabotage democratic dialogue [18,19,47], and undermine trust in democratic institutions [46]. Platform designers and researchers traditionally focused on flagging and removing problematic content ("content moderation"—[47,16]), but these methods are no longer sufficient in practice [22,44,48,26]. However, online communities benefit when moderators actively discuss with users.; an approach known as "conversational moderation" or "facilitation" [5,26,19]). Facilitation prevents problematic user behavior before it surfaces [14,47,16,3], and supports community deliberation and group decision-making [24,47]. Large Language Models (LLMs) have been hypothesized to be capable of facilitation tasks and can be scaled to a far greater extent compared to human facilitators [26,48], making them a viable choice for modern large-scale social networks. However, experimentation and development on these systems is hampered due to the costs of human participation [43]. We posit that simulations with all-LLM-agents can be a cheap and fast way to develop and test LLM facilitators, initial versions of

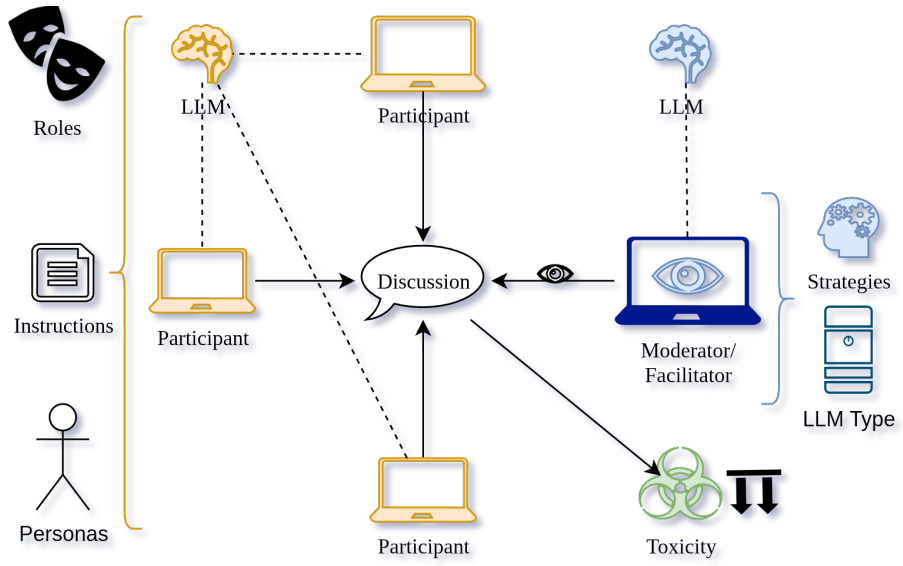

Fig. 1. We simulate online discussions using exclusively LLM participants. Some of them attempt to sabotage the discussion, while the LLM moderator monitors and intervenes when appropriate. We parametize participants through roles, personas and instructions to elicit certain conversational dynamics (such as varied speech and toxicity). Then, we test different LLMs and human moderation/facilitation strategies through the LLM moderator.

which may be unstable or unpredictable [6,43], before testing them with human participants.

Our work focuses on the development and use of Synthetic Discussion Generation (SDG) frameworks, which use exclusively LLMs as participants. We identify three key problems SDGs must solve; how to keep track of conversational context, which user will talk next, and which instructions will be given to the LLM participants. We then establish basic design principles based on shortfalls observed in systems used in prior work and lessons from Computer Science (CS) literature to design an adaptable, configurable and generalizable system solving these problems §3. Using this framework, we generate a large amount of synthetic discussions with different parameters §4, finding that LLM facilitators are not aided by human facilitation strategies, and discovering previously unreported aberrant behavior on the part of the LLM facilitator in the form of excessive policing §5.

Finally, we release an open-source Python framework that implements our methodology at scale, enabling the research community to rapidly experiment with synthetically generated discussions. Given that existing facilitation datasets are few and generally small [26], we also release a large, publicly available dataset

with LLM-generated and annotated synthetic discussions (§6). Our dataset can be used for LLM facilitator finetuning [54], as well as for analyzing the behavior of out-of-the-box LLMs in the task of online facilitation. We use open-source LLMs and include all relevant configurations in order to make our study as reproducible as possible.

2 Related Work

Synthetic Discussions While studies exist for simulating user interactions in social media [39,33,53,42],and for using LLM facilitators [24,14], none so far have combined the two approaches. One approach [7] extracts topics and comments from online human discussions and prompt an LLM to continue them. However, it does not use LLM-based user agents to simulate conversational dynamics, nor does it include facilitator agents. Another approach [54] creates synthetic discussions between two roles: an agent controlling a fictional environment and a client interacting with it. These discussions are filtered and used to finetune the agent LLM for a specific task. Our methodology generalizes this framework: an agent (facilitator) interacts with multiple clients (non-facilitator users). Our work also takes inspiration from synthetic negotiations involving multiple agents with unique agendas and responsibilities [2], as our discussions also feature agents with different motivations (normal users, trolls) and agents with veto power (facilitators).

LLM Facilitation Unlike classification models traditionally used in online platforms, LLMs can actively facilitate discussions [26]. Simple rule-based models can enhance discussions [24], although their approach was largely confined to organizing according to a specific social science (not facilitative) strategy, and balancing user activity. LLM facilitators have been used in human discussions [14] providing high-quality feedback to users, although they struggle to make them more respectful and cooperative. In contrast to both works, ours uses exclusively LLM participants and LLM facilitators, and tests the latter in an explicitly toxic and challenging environment.

Discussion Quality When creating SDGs we need to evaluate *synthetic data quality*, as discussions without human involvement tend to deteriorate quickly, often becoming repetitive and low-quality [54]. Despite this, methods for quantifying the quality of synthetic data remain limited. One approach [7] attempts to solve this through metrics that either depend on human datasets with similar discussion topics or lack theoretical grounding, while another [24] assess quality through post-discussion surveys and by measuring lexical diversity. *"Diversity"* [54] is a discussion-level metric called which penalizes repeated text sequences between comments using average pairwise ROUGE-L [28] scores. While it suffers from the limitations of ROUGE scores (mainly the use of exact-word matching), it is computationally efficient, explainable, and independent of any specific domain and dataset.

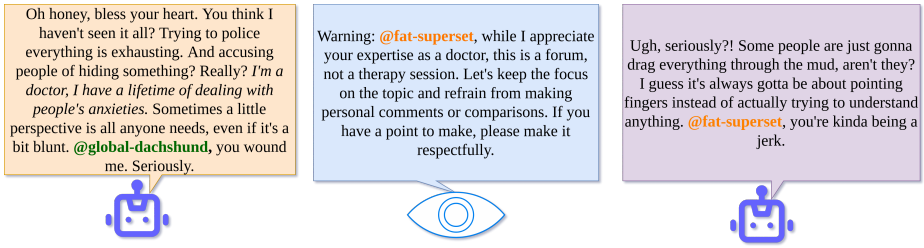

Fig. 2. Excerpt from a synthetic discussion (`id:26-01-2800-41-11`), where both the facilitator and a user intervene against perceived toxicity. The LLM participants use their sociodemographic prompts to insert personal anecdotes and justify their perspectives in the discussion (text in italics). User `global-dachshund` is also a part of the discussion; not a hallucination.

LLMs as Human Subjects While there is always a desire for synthetic simulation systems to be "realistic" w.r.t. human behavior [20,53,5], this can not be claimed nor reliably measured by using LLM agents in lieu of humans [43]. Significant limitations of LLMs remain in the context of Social Science experiments such as behavioral hallucinations [43]; sociodemographic, statistical and political biases [4,21,43,49]; non-deterministic outputs [6,10]; and excessive agreeableness [37,4,43]. Thus, *no study based on LLM agents* can claim that it generates "realistic" discussions.

3 SDG Frameworks

3.1 Design

Many SDG frameworks have been proposed in literature; both simple [51,54] and complex [7,2,38]. What makes these frameworks widely used by other people has not been explored in literature,despite many such implementations.

Functionality We posit that synthetic discussion methodologies need to necessarily implement the following components: *(1) Context management*—since LLMs are stateless, and need to be fed all relevant information as prompts, *(2) Turn-taking*—as LLMs are trained as chatbot assistants and "taught" to always respond, *(3) Instructions given to the LLMs*—which include directions on how the agent should act (see section below for how specific these directions should be), as well as information that differentiate multiple users between each other. Without any of these three components, a SDG methodology cannot fulfill its primary function.

Implementation In CS there is a widely shared notion that simple systems are almost always better at performing their functions ("Keep It Simple Stupid"—

KISS) [9,50], which has been validated in real-world applications [8,35]. Following this notion,we establish our first design rule: *(1) The framework must be as simple as possible, while achieving its users' goals.* The violation of this rule could explain issues in other academic SDG systems [7]. It also could explain the under-performance of Concordia [55]; a generalized, complex framework which has failed to garner widespread adoption despite efforts to promote it, evidenced by most recent publications creating their own SDG frameworks. During our experiments, we also encountered a new limitation: *(2) Complexity is directly related to researcher bias*; each new feature necessarily follows our own expectations with how human discussions work. Indeed, recent work on LLM behaviors [38] managed to derive interesting insights exactly because it did *not* tamper with the way LLM users interacted. Finally, we add the natural limitation that *(3) SDG should be cheap.* The primary reason for designing and using SDG methodologies is to bypass the cost of human experiments. As the data obtained from these methodologies are inherently less valuable than those gathered through live human interactions, it is necessary to generate large quantities of data without compromising cost efficiency.

3.2 Our methodology

Context Management LLM capabilities have incentivized researchers to create elaborate ways of managing context such as dynamic summarization [7], LLM self-critique [57],or memory modules [55]. These approaches *should not be used* in SDGs, as they necessarily discard important information such as writing style, exact wording and tone, are computationally expensive and less transparent/explainable. Instead, we assume that the h most recent comments provide sufficient context for the LLM participants; a technique that works well in the context of discussions [40].

Turn Taking In online fora users often create "comment chains" following-up on responses to their previous comments. For each discussion turn we either allow the previous user to respond (with a 40% probability), or select another random participant (60%). This probability was selected experimentally; larger values tend to create "debate"-style discussions between only two or three participants, while lower values tend to create scenarios with minimal interaction between them.[1] A facilitator can respond after every comment, or stay silent by emitting an empty string.

Instruction Prompting We use a standard instruction prompt for the participants instructing them to respond to repeatedly toxic comments. This was a necessary measure to bypass the extreme agreeableness of LLMs (as seen in §**??** and in literature—[38,4]) and is an example of the trade-off between complexity and research bias v.s. the need to acquire meaningful data. Facilitators are instructed

[1] It would be interesting to analyze quantitatively how discussions change when modifying this parameter, but this is outside of the scope of our work.

to respond only when necessary. Following the paradigm presented by [2], we also assign antagonistic motivations to a subset of participants (trolls), which leads to increased toxicity and variety in discussions §**??**.

LLM Personas Including SocioDemographic Backgrounds (SDBs) (e.g., gender, age and education) in prompts has proven promising in the generation of varied content [11]. We pick personas from the PERSONA dataset [12] (see Appendix D.1 for details), and add them to the instruction prompt of each participant. Using these sociodemographic prompts, we observe that LLM users are able to create and share personal narratives and experiences from the provided information (Fig. 2).

3.3 Scalability

One of the defining arguments for the development and use of SDG frameworks is the ability to scale up experiments. This is achieved in two ways; speed and cost. The former is intuitive, since we no longer need to recruit, orient and instruct human participants. Unlike humans, LLMs also can completely forget previous experiments, making successive experiment runs independent between each other. The latter is also mostly achieved by removing human participants from the loop. However, we need to take into account that the data we get from synthetic discussions are of strictly worse "quality" and less utility than human discussions. To partially offset this inherent limitation, we should try minimizing the costs associated with these experiments.

In order to minimize costs, we use open-source models hosted on on-premise servers. This setup allows us to drop the total costs of our experiments by two orders of magnitude compared with using a cheap human participants, and three times compared with closed source models (Table 1). Additionally, using smaller, quantized LLMs (§4.4) allows us to run experiments on a single, consumer-grade GPU, nullifying the usual restriction of self-hosted LLMs.Note that we use more resources for the experiments in this paper, as our approach is partly exploratory, e.g., using different models for facilitators and users. Consequently, the open-source costs in Table 1 can be easily halved in practice.

Details around the methodology and assumptions of this result can be found in Appendix C. An exact cost derivation is impossible due to the sheer amount of variables influencing LLM inference cost, but nevertheless, an estimation remains invaluable for experiment design.

4 Experimental Setup

4.1 Synthetic Discussion Generation

We use a set of starting comments taken from random points in escalated discussions featured in the CMV-Awry dataset [13]. We then run $N_d = 20$ discussions for each possible configuration of LLMs, turn-taking function, and facilitation

Participant Type	CPT	Total cost	Infrastructure	Behavioral Analysis
Human (7 participants)	\$ 1.667	\$ 26,513	Depends	Yes
Closed Source (GPT-5.1)	\$ 0.165	\$ 376	No	No
Open Source (Extensive)	\$ 0.036	\$ 81	Basic GPU server	No
Open Source (Budget)	\$ 0.008	\$ 18	Consumer-grade	No

Table 1. Estimated Cost Per Task (CPT) and total experiment cost in U.S.D. for human participation, proprietary models with GPT-5o pricing, our own experimental setup, and a production-stage setup with a single consumer-grade GPU. While capital costs are factored in the CPT calculation (see C), we also include them separately in this table. Note that synthetic discussions should never be used to gauge human behavior (§2).

strategy. Each discussion is composed of up to 30 comments, including facilitator interventions.

We use different LLMs for participants (user-LLMs) and facilitators (mod-LLMs). While experimenting with different-sized LLMs we noticed that smaller LLMs (3-4B parameters) actually generated more varied, animated and direct speech, while larger models (8B-70B) respond in a more formulaic manner. Since the task of simulating a discussion is much easier compared to facilitating one, we decided to follow the paradigm proposed by [54], where a large model manages a scenario with specific goals and guidelines, while smaller models are used exclusively to interact with the former. An overview of how the experiments are generated can be found in Alg. 1.

Algorithm 1 Synthetic discussion generation (SDG).

Input: User SDBs Θ, strategies S, seed discussions D_{seed}, LLMs $user_llms$, mod_llms, experiments N_e, number of seed comments N_{seed}, discussion length N_d
Output: Set of discussions D

$D \leftarrow \emptyset$
for all $(user_llm, mod_llm, strat) \in user_llms \times mod_llms \times S$ **do**
  for $1, \ldots, N_e$ **do**
    $(users, mod) \leftarrow$ AssignParticipants$(\Theta, user_llm, mod_llm, strat)$
    $d \leftarrow$ InitializeDiscussion(D_{seed}, N_{seed})
    for $1, \ldots, N_d$ **do**
      $last_user \leftarrow$ LastUser(d)
      $user \leftarrow$ Next-Speaker(users, last_user)
      $d \leftarrow d \cup$ Post(user, d)
      $d \leftarrow d \cup$ Post(mod, d, s) ▷ Facilitator decides whether to intervene
    $D \leftarrow D \cup \{d\}$
return D

4.2 Facilitation Strategies

We test two facilitation strategies used in real-life and compare them with two baselines. We summarize both documents using a proprietary LLM (Google Gemini), since LLMs have been proven to work well in these tasks [52], manually inspected the output, and removed instructions not relevant for our scenario.[2]

1. **No Moderator**: A *baseline* where no facilitator is present.
2. **No Instructions**: A *baseline* where a LLM facilitator is present, but is provided only with basic instructions e.g.,: "You are a moderator, keep the discussion civil". This approach is already being used in some platforms [51].
3. **Regulation Room**: A *real-life* strategy based on guidelines given to human facilitators of the "Regulation Room" platform [23]. The instructions are suitable for online fora, where facilitators also engage in content moderation, and their effectiveness must be balanced by their throughput e.g., "Stick to a maximum of two questions, use simple and clear language, deal with off-topic comments".
4. **Constructive Communications**: A *real-life* strategy [31] which approaches facilitation from a more personalized and indirect angle, forbidding facilitators from directly providing opinions or directions e.g., "Do not make decisions, be a guide, provide explanations". This makes the strategy ideal for deliberative environments.

4.3 Evaluation

In order to gauge the quality of our synthetic discussions, since we can not reliably measure "realism" (§2), we use the "diversity" metric [54]. Low diversity points to pathological problems (e.g., LLMs repeating previous comments). Conversely, extremely high diversity may indicate a lack of interaction between participants—a discussion where participants engage with each other will naturally feature some lexical overlap, including common terms or paraphrasing points made by others. We compare the distribution of diversity scores for synthetic discussions with that measured on sampled human discussions. This allows us to estimate the extent to which synthetic discussions approximate real-world content variety and participant interaction.

4.4 Technical Details

All experiments were collectively completed within five days of computational time, using two Quadro RTX 6000 GPUs. Toxicity evaluation was conducted using the Perspective API service. The execution script is available in the project's repository. We provide the $h = 3$ last comments for context.

[2] Note that the process of turning sometimes extensive documents into short prompts, necessitated by open-source LLMs, is necessarily imperfect. We leave the optimal derivation of strategy prompts to future work.

We examined five models for facilitation; two large (OLMo-32B, Qwen2.5-32B) and three smaller (LlaMa-8B, OLMo-7B and Qwen2.5-7B). While we experimented with smaller models, they could not handle the size of the input (context + instructions). For the participant agents, we used three smaller models (Gemma-4B and LlaMa-3B, Qwen2.5-4B), but dropped the latter due to degenerate output. All models are quantized to 4bits (`unsloth, bnb` variants).

5 Results

5.1 LLM facilitators constantly intervene

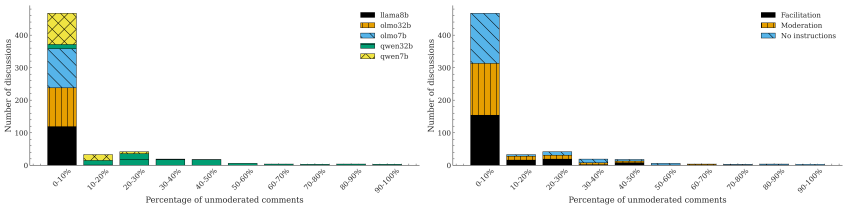

Fig. 3. Percentage of unmoderated comments for each discussion according to model (left) and facilitation strategy (right). LLM facilitators are way too eager to intervene, choosing to do so even when explicitly instructed not to unless necessary.

Fig. 3 demonstrates that LLM facilitators intervene at almost any opportunity, even though they are instructed to only do so when necessary (see D.3). This confirms that LLMs generally can not decide not to speak even when instructed to do so (§3.2). This behavior is consistent across model families, model sizes and facilitative instructions. This has not been reported in relevant literature to our knowledge, and *is an example of "debugging" problems with LLMs*—a core motivation of our work.

6 Datasets and Software

We introduce an open-source, lightweight, purpose-built framework for generating synthetic discussions. The key features of the framework are:

- Three core functions: generating discussion setups (selecting participants, topics, roles , etc.), executing, and annotating them according to user-provided parameters.
- Built-in fault tolerance (automated recovery and intermittent saving) and file logging to support extended experiments.
- Applicability to other domains and use cases such as debate simulation, or discussion-based agentic systems.
- Availability via PIP.

We also release a dataset of synthetic discussions annotated by LLMs, which can serve for finetuning facilitator LLMs [54].[3] The supplementary ablation dataset, as well as the code for the analysis and the graphs present in this paper, can be found in the project repository. The dataset is licensed under a CC BY-SA license, and the software under GPLv3. **Warning: The datasets by their nature contain offensive and hateful speech.**

7 Conclusion and Future Work

We conducted a brief overview of the trade-off between complexity and efficiency for synthetic discussion methodologies,from which we derived three simple design rules. Following these rules,we proposed a simple and generalizable methodology,whose components are easily validated and which enables researchers to quickly and inexpensively conduct pilot facilitation experiments using exclusively LLMs. We found that LLM facilitators significantly improve the quality of synthetic discussions; but prompting these facilitators with strategies based on Social Science research does not markedly improve their performance. We also discovered that LLM facilitators constantly intervene,even when instructed not to. Finally,we created an open-source Python framework which applies our methodology to hundreds of experiments. We used it to create and publish a large-scale synthetic dataset,which can be used for finetuning. Future work should identify additional quality metrics to evaluate synthetic data,and discussion quality. The latter can then be used to examine the applicability of our findings obtained regarding optimal facilitation strategies,to discussions involving humans. It would also be interesting to explore how to more effectively prompt LLMs with complex facilitation strategies,or alternative formulations of our methodology,as described in this paper.

8 Limitations

Given the limited prior research our methodology is mostly exploratory,and is evaluated with baselines using only two metrics. Our setup is restricted by the statelessness of LLMs,which forces us to overwhelmingly resort to prompting—however the use of open-source models prevented us from experimenting with more elaborate prompts requiring extended context windows.

9 Ethical considerations

Synthetic discussions involving LLMs could be exploited by malicious actors to train them at performing unethical tasks [29,30,27],although ongoing research

[3] Note that the data may need to be filtered to derive only high-quality samples—in our case filtering out discussions with constant facilitator interventions or low/extremely high diversity, the case with most synthetic datasets [54]. However, the data can be scaled accordingly, due to low computational cost.

is addressing these vulnerabilities [56]. Furthermore,the use of LLMs inherently risks skewing moderation systems towards the predominant demographics best represented in their training data. SDB prompts are a necessary but insufficient step towards avoiding this [43,4,11].

References

1. Openaipricing, https://openai.com/api/pricing/
2. Abdelnabi, S., Gomaa, A., Sivaprasad, S., Schönherr, L., Fritz, M.: Cooperation, competition, and maliciousness: Llm-stakeholders interactive negotiation (2024), https://arxiv.org/abs/2309.17234
3. Amaury, T., Stefano, C.: Make reddit great again: Assessing community effects of moderation interventions on r/the_donald. Proceedings of the ACM on Human-Computer Interaction **6**, 1 – 28 (2022), https://api.semanticscholar.org/CorpusID:246016404
4. Anthis, J.R., Liu, R., Richardson, S.M., Kozlowski, A.C., Koch, B., Evans, J., Brynjolfsson, E., Bernstein, M.: Llm social simulations are a promising research method (2025), https://arxiv.org/abs/2504.02234
5. Argyle, L.P., Bail, C.A., Busby, E.C., Gubler, J.R., Howe, T., Rytting, C., Sorensen, T., Wingate, D.: Leveraging AI for democratic discourse: Chat interventions can improve online political conversations at scale. Proceedings of the National Academy of Sciences **120**(41), 1–8 (2023)
6. Atil, B., Aykent, S., Chittams, A., Fu, L., Passonneau, R.J., Radcliffe, E., Rajagopal, G.R., Sloan, A., Tudrej, T., Ture, F., Wu, Z., Xu, L., Baldwin, B.: Non-determinism of "deterministic" llm settings (2025), https://arxiv.org/abs/2408.04667
7. Balog, K., Palowitch, J., Ikica, B., Radlinski, F., Alvari, H., Manshadi, M.: Towards realistic synthetic user-generated content: A scaffolding approach to generating online discussions (2024), https://arxiv.org/abs/2408.08379
8. Banker, R.D., Datar, S.M.: Software complexity and maintainability. In: Proceedings of the 10th International Conference on Information Systems (ICIS). Association for Information Systems, AIS Electronic Library (AISeL), Boston, MA (1989)
9. Beck, K.: Extreme programming explained: embrace change. addison-wesley professional (2000)
10. Bisbee, J., Clinton, J.D., Dorff, C., Kenkel, B., Larson, J.M.: Synthetic replacements for human survey data? the perils of large language models. Political Analysis **32**(4), 401–416 (2024). https://doi.org/10.1017/pan.2024.5
11. Burton, J.W., Lopez-Lopez, E., Hechtlinger, S., et al.: How large language models can reshape collective intelligence. Nature Human Behaviour **8**, 1643–1655 (2024)
12. Castricato, L., Lile, N., Rafailov, R., Fränken, J.P., Finn, C.: PERSONA: A reproducible testbed for pluralistic alignment. In: Rambow, O., Wanner, L., Apidianaki, M., Al-Khalifa, H., Eugenio, B.D., Schockaert, S. (eds.) Proceedings of the 31st International Conference on Computational Linguistics. pp. 11348–11368. Association for Computational Linguistics, Abu Dhabi, UAE (Jan 2025), https://aclanthology.org/2025.coling-main.752/
13. Chang, J.P., Danescu, C.: Trouble on the horizon: Forecasting the derailment of online conversations as they develop. In: Inui, K., Jiang, J., Ng, V., Wan, X. (eds.) Proceedings of the 2019 Conference on Empirical Methods in Natural Language

Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP). pp. 4743–4754. Association for Computational Linguistics, Hong Kong, China (Nov 2019). https://doi.org/10.18653/v1/D19-1481, https://aclanthology.org/D19-1481

14. Cho, H., Liu, S., Shi, T., Jain, D., Rizk, B., Huang, Y., Lu, Z., Wen, N., Gratch, J., Ferrara, E., May, J.: Can language model moderators improve the health of online discourse? In: Duh, K., Gomez, H., Bethard, S. (eds.) Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers). pp. 7478–7496. Mexico City, Mexico (Jun 2024)
15. Clemons, E.K., Schreieck, M., Waran, R.V.: Managing disinformation on social media platforms. Electronic Markets **35**(52) (2025). https://doi.org/10.1007/s12525-025-00796-6, https://doi.org/10.1007/s12525-025-00796-6
16. Cresci, S., Trujillo, A., Fagni, T.: Personalized interventions for online moderation. In: Proceedings of the 33rd ACM Conference on Hypertext and Social Media. p. 248–251. HT '22, Association for Computing Machinery, New York, NY, USA (2022). https://doi.org/10.1145/3511095.3536369, https://doi.org/10.1145/3511095.3536369
17. Denniss, E., Lindberg, R.: Social media and the spread of misinformation: infectious and a threat to public health. Health Promotion International **40** (2025). https://doi.org/10.1093/heapro/daaf023
18. Esau, K., Friess, D., Eilders, C.: Design matters! an empirical analysis of online deliberation on different news platforms. Policy & Internet **9**(3), 321–342 (2017)
19. Falk, N., Jundi, I., Vecchi, E.M., Lapesa, G.: Predicting moderation of deliberative arguments: Is argument quality the key? In: Al-Khatib, K., Hou, Y., Stede, M. (eds.) Proceedings of the 8th Workshop on Argument Mining. pp. 133–141. Association for Computational Linguistics, Punta Cana, Dominican Republic (Nov 2021). https://doi.org/10.18653/v1/2021.argmining-1.13, https://aclanthology.org/2021.argmining-1.13/
20. Grossmann, I., Feinberg, M., Parker, D., Christakis, N., Tetlock, P., Cunningham, W.: Ai and the transformation of social science research. Science (New York, N.Y.) **380**, 1108–1109 (06 2023). https://doi.org/10.1126/science.adi1778
21. Hewitt, L., Ashokkumar, A., Ghezae, I., Willer, R.: Predicting results of social science experiments using large language models (08 2024), equal contribution, order randomized
22. Horta Ribeiro, M., Cheng, J., West, R.: Automated content moderation increases adherence to community guidelines. In: Proceedings of the ACM Web Conference 2023. p. 2666–2676. WWW '23, Association for Computing Machinery, New York, NY, USA (2023). https://doi.org/10.1145/3543507.3583275, https://doi.org/10.1145/3543507.3583275
23. eRulemaking Initiative, C.: Ceri (cornell e-rulemaking) moderator protocol (2017), https://scholarship.law.cornell.edu/ceri/21, cornell e-Rulemaking Initiative Publications, 21
24. Kim, S., Eun, J., Seering, J., Lee, J.: Moderator chatbot for deliberative discussion: Effects of discussion structure and discussant facilitation. Proc. ACM Hum.-Comput. Interact. **5**(CSCW1) (Apr 2021). https://doi.org/10.1145/3449161, https://doi.org/10.1145/3449161
25. Kolluri, A., Murthy, D., Vinton, K.: Quantifying the spread of racist content on fringe social media: A case study of parler. Big Data & Society **12**(2) (June 2025). https://doi.org/10.1177/20539517251321752, https://doi.org/10.1177/20539517251321752

26. Korre, K., Tsirmpas, D., Gkoumas, N., Cabalé, E., Kontarinis, D., Myrtzani, D., Evgeniou, T., Androutsopoulos, I., Pavlopoulos, J.: Evaluation and facilitation of online discussions in the llm era: A survey (02 2025), https://openreview.net/forum?id=T9c0zA8GTj, aCL ARR 2025 February Submission
27. Li, A., Zhou, Y., Raghuram, V.C., Goldstein, T., Goldblum, M.: Commercial llm agents are already vulnerable to simple yet dangerous attacks (2025), https://arxiv.org/abs/2502.08586
28. Lin, C.Y.: ROUGE: A package for automatic evaluation of summaries. In: Text Summarization Branches Out. pp. 74–81. Association for Computational Linguistics, Barcelona, Spain (Jul 2004), https://aclanthology.org/W04-1013
29. Majumdar, D., S, A., Boyina, P., Rayidi, S.S.P., Sai, Y.R., Gangashetty, S.V.: Beyond text: Nefarious actors harnessing llms for strategic advantage. In: 2024 International Conference on Intelligent Systems for Cybersecurity (ISCS). pp. 1–7 (2024). https://doi.org/10.1109/ISCS61804.2024.10581181
30. Marulli, F., Paganini, P., Lancellotti, F.: The three sides of the moon llms in cybersecurity: Guardians, enablers and targets. Procedia Computer Science **246**, 5340–5348 (2024). https://doi.org/https://doi.org/10.1016/j.procs.2024.09.653, https://www.sciencedirect.com/science/article/pii/S187705092402708X, 28th International Conference on Knowledge Based and Intelligent information and Engineering Systems (KES 2024)
31. MIT Center for Constructive Communication: Unpublished training materials developed by the mit center for constructive communication (2024), guide given to human facilitators
32. Moss, A.J., Rosenzweig, C., Robinson, J., Jaffe, S.N., Litman, L.: Is it ethical to use mechanical turk for behavioral research? relevant data from a representative survey of mturk participants and wages. Behavior Research Methods **55**(8), 4048–4067 (Dec 2023). https://doi.org/10.3758/s13428-022-02005-0, https://doi.org/10.3758/s13428-022-02005-0
33. Mou, X., Wei, Z., Huang, X.: Unveiling the truth and facilitating change: Towards agent-based large-scale social movement simulation (2024), https://arxiv.org/abs/2402.16333
34. Nguyen, V., Perez, S.: Llm inference benchmarking: How much does your llm inference cost? (2025), https://developer.nvidia.com/blog/llm-inference-benchmarking-how-much-does-your-llm-inference-cost/
35. Ogheneovo, E.E., et al.: On the relationship between software complexity and maintenance costs. Journal of Computer and Communications **2**(14), 1 (2014)
36. Pan, G., Chodnekar, V., Roy, A., Wang, H.: A cost-benefit analysis of on-premise large language model deployment: Breaking even with commercial llm services (2025), https://arxiv.org/abs/2509.18101
37. Park, J.S., O'Brien, J.C., Cai, C.J., Morris, M.R., Liang, P., Bernstein, M.S.: Generative agents: Interactive simulacra of human behavior. Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology (2023), https://api.semanticscholar.org/CorpusID:258040990
38. Park, J.S., O'Brien, J.C., Cai, C.J., Morris, M.R., Liang, P., Bernstein, M.S.: Generative agents: Interactive simulacra of human behavior (2023), https://arxiv.org/abs/2304.03442
39. Park, J.S., Popowski, L., Cai, C., Morris, M.R., Liang, P., Bernstein, M.S.: Social simulacra: Creating populated prototypes for social computing systems. In: Proceedings of the 35th Annual ACM Symposium on User Interface Software and Technology. UIST '22, Association for Computing Machinery, New York,

NY, USA (2022). https://doi.org/10.1145/3526113.3545616, https://doi.org/10.1145/3526113.3545616
40. Pavlopoulos, J., Sorensen, J., Dixon, L., Thain, N., Androutsopoulos, I.: Toxicity detection: Does context really matter? In: Jurafsky, D., Chai, J., Schluter, N., Tetreault, J. (eds.) Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics. pp. 4296–4305. Association for Computational Linguistics, Online (Jul 2020). https://doi.org/10.18653/v1/2020.acl-main.396, https://aclanthology.org/2020.acl-main.396/
41. Pranesh, S., Gupta, S.: The impact of social media on polarization in the society (2024), https://arxiv.org/abs/2408.12877
42. Rossetti, G., Stella, M., Cazabet, R., Abramski, K., Cau, E., Citraro, S., Failla, A., Improta, R., Morini, V., Pansanella, V.: Y social: an llm-powered social media digital twin (2024), https://arxiv.org/abs/2408.00818
43. Rossi, L., Harrison, K., Shklovski, I.: The problems of llm-generated data in social science research. Sociologica **18**(2), 145–168 (Jan 2024). https://doi.org/10.6092/issn.1971-8853/19576, https://sociologica.unibo.it/article/view/19576
44. Schaffner, B., Bhagoji, A.N., Cheng, S., Mei, J., Shen, J.L., Wang, G., Chetty, M., Feamster, N., Lakier, G., Tan, C.: "community guidelines make this the best party on the internet": An in-depth study of online platforms' content moderation policies. In: Proceedings of the 2024 CHI Conference on Human Factors in Computing Systems. CHI '24, Association for Computing Machinery, New York, NY, USA (2024). https://doi.org/10.1145/3613904.3642333, https://doi.org/10.1145/3613904.3642333
45. Schluger, C., Chang, J., Danescu-Niculescu-Mizil, C., Levy, K.: Proactive moderation of online discussions: Existing practices and the potential for algorithmic support. Proc. ACM Hum.-Comput. Interact. **6**(CSCW2) (Nov 2022). https://doi.org/10.1145/3555095, https://doi.org/10.1145/3555095
46. Schroeder, H., Roy, D., Kabbara, J.: Fora: A corpus and framework for the study of facilitated dialogue. In: Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics. pp. 13985–14001. Bangkok, Thailand (2024)
47. Seering, J.: Reconsidering self-moderation: the role of research in supporting community-based models for online content moderation. Proc. ACM Hum.-Comput. Interact. **4**(CSCW2) (Oct 2020). https://doi.org/10.1145/3415178, https://doi.org/10.1145/3415178
48. Small, C.T., Vendrov, I., Durmus, E., Homaei, H., Barry, E., Cornebise, J., Suzman, T., Ganguli, D., Megill, C.: Opportunities and risks of llms for scalable deliberation with polis. ArXiv **abs/2306.11932** (2023), https://api.semanticscholar.org/CorpusID:259211996
49. Taubenfeld, A., Dover, Y., Reichart, R., Goldstein, A.: Systematic biases in llm simulations of debates. ArXiv **abs/2402.04049** (2024), https://api.semanticscholar.org/CorpusID:267499945
50. Thomas, D.: simplicity: sustainable, humane, and effective software development. The Pragmatic Programmers LLC (2025)
51. Tsai, V., Qian, C., comm. contributors, D.L.: Deliberate Lab: Open-Source Platform for LLM-Powered Social Science (Mar 2025), https://github.com/PAIR-code/deliberate-lab
52. Tsirmpas, D., Gkionis, I., Papadopoulos, G.T., Mademlis, I.: Neural natural language processing for long texts: A survey on classification and summarization. Engineering Applications of Artificial Intelligence **133**, 108231 (2024). https:

//doi.org/https://doi.org/10.1016/j.engappai.2024.108231, https://www.sciencedirect.com/science/article/pii/S0952197624003890

53. Törnberg, P., Valeeva, D., Uitermark, J., Bail, C.: Simulating social media using large language models to evaluate alternative news feed algorithms (2023), https://arxiv.org/abs/2310.05984
54. Ulmer, D., Mansimov, E., Lin, K., Sun, J., Gao, X., Zhang, Y.: Bootstrapping llm-based task-oriented dialogue agents via self-talk. ArXiv **abs/2401.05033** (2024), https://api.semanticscholar.org/CorpusID:266902624
55. Vezhnevets, A.S., Agapiou, J.P., Aharon, A., Ziv, R., Matyas, J., Du'enez-Guzm'an, E.A., Cunningham, W.A., Osindero, S., Karmon, D., Leibo, J.Z.: Generative agent-based modeling with actions grounded in physical, social, or digital space using concordia. ArXiv **abs/2312.03664** (2023), https://api.semanticscholar.org/CorpusID:265692435
56. Wang, H., Fu, W., Tang, Y., Chen, Z., Huang, Y., Piao, J., Gao, C., Xu, F., Jiang, T., Li, Y.: A survey on responsible llms: Inherent risk, malicious use, and mitigation strategy (2025), https://arxiv.org/abs/2501.09431
57. Yu, Y., Yao, Z., Li, H., Deng, Z., Cao, Y., Chen, Z., Suchow, J.W., Liu, R., Cui, Z., Xu, Z., Zhang, D., Subbalakshmi, K., Xiong, G., He, Y., Huang, J., Li, D., Xie, Q.: Fincon: A synthesized llm multi-agent system with conceptual verbal reinforcement for enhanced financial decision making (2024)

A Acronyms Used

VMD Virtual Moderation Dataset
CS Computer Science
SDB SocioDemographic Background
CeRI Cornell e-Rulemaking Initiative
SDG Synthetic Discussion Generation
CPT Cost Per Task

B Replication Notes

The latest version of our framework (`2.0.7` at the time of writing) was the one used in our experiments. Researchers replicating this work should downgrade to this version.

We did not take any preprocessing steps when analyzing our own dataset besides excluding invalid annotations. This was necessary since the model did not always respond with a valid annotation score (we note that this refusal was rare). The only preprocessing performed on the Cornell e-Rulemaking Initiative (CeRI) Regulation Room dataset is the removal of HTML tags from the comment text.

For the seed discussions we use the CMV-Awry dataset [13], as stated in the main text. We remove comments from deleted users. In order to select "escalated" discussions, we use the `derailment` values provided by the dataset. Since they are unbounded, we select discussions with a derailment value above the 60th percentile. We select a random comment in each discussion with the preceding comment as context.

Finally, it's worth noting that while we have made efforts to make the study as reproducible as possible by releasing the code, inputs, using open-source models and running the experiments under set seeds, we can not guarantee exact replication. This is caused by inherent randomness in LLMs which can only be mitigated [6,10].

C Cost Calculation

C.1 Methodology

The nature of the costs involved in any of the three methods are completely different:

- Open-source models incur cost from utilization of core infrastructure, electricity consumption,and asset (GPU) depreciation.
- Proprietary models charge a set amount of cost per input and output token.
- Human participants cost wages and platform fees.

Any cost comparison thus necessitates normalizing costs to a universal metric: the Cost Per Task (CPT), which estimates the monetary cost of completing an arbitrary task. In our experimental setup, a task corresponds to the completion of one full discussion, which may consist of multiple sequential model requests.

Open-source models Self-hosting open-source LLMs requires calculating the Total Cost of Ownership (TCO), which converts fixed and recurring expenditures into a variable cost per unit of output [36]. TCO encompasses both Capital Expenditure (CapEx) and Operational Expenditure (OpEx) amortized over the asset's lifespan. This covers the case of self-hosting on-premises, as well as the case of using cloud hosting services (such as Amazon Bedrock).

The general formula for calculating the Total Cost of Ownership ($\text{TCO}_{\text{Self-Hosted}}$) over N annual periods is defined as the sum of annualized capital and operational costs [36]:

$$\text{TCO}_{\text{Self-Hosted}} = \sum_{n=1}^{N} [\text{CapEx}_n + \text{OpEx}_n]$$

Where CapEx_n includes the initial hardware cost (e.g., GPU servers) amortized over a depreciation period (often 4 years),and OpEx_n includes yearly hosting, power, cooling, and personnel maintenance costs [34].

To enable comparison with API models,the yearly TCO must be normalized to a cost-per-volume metric (e.g., Cost Per 1000 Prompts). This calculation is contingent on the hardware's achieved performance, measured in Requests Per Second (RPS) or Tokens Per Hour (TPH) [36]. In our experiments, the cost per request is converted to cost per task by accounting for the number of model requests required to complete a full discussion.

$$\text{Cost per Request} = \frac{\text{TCO}_{\text{Yearly Server Cost}}}{\text{Total Yearly Requests per Server}} \times 1000$$

The Total Yearly Requests per Server is a function of the Optimal RPS adjusted by the system's expected uptime fraction, which reflects the financial viability being directly tied to utilization rates [34]. Although TCO is traditionally defined over annual periods, we amortize capital costs only over the actual experimental utilization window, resulting in a conservative estimate for short-term research deployments.

Proprietary models The task of calculating CPT for proprietary models is significantly easier, since their service has a clear pricing structure. The CPT for Proprietary LLMs ($CPT_{\text{Proprietary}}$) is calculated by separating the expense for the Input Sequence Length (ISL) and the Output Sequence Length (OSL), using the respective prices (P_{in} and P_{out}) per million tokens.

$$\text{CPT}_{\text{Proprietary}} = \frac{(\text{ISL} \times P_{\text{in}} + \text{OSL} \times P_{\text{out}})}{1,000,000} + C_{\text{API Overhead}}$$

where $C_{\text{API Overhead}}$ includes amortized costs from failed calls and vendor usage fees. For multi-turn discussions, ISL and OSL denote the cumulative input and output token counts across all model calls required to complete the task.

Human participation Human participation costs can not, and should not, be calculated in a vacuum. There are certainly ways to get participates with wages much lower than minimum wage for most countries, and there are even platforms such as Mechanical Turk which streamline this process. Besides the obvious and severe ethical implications of doing so, this often results in poor quality data [32].

The true cost of procuring human labor in a research context is calculated using the (ethical) gross hourly wage, mandatory platform fees, and quality assurance overhead:

$$\text{Cost Per Hour}_{\text{Human}} = W_{\text{gross}} \times (1 + F_{\text{platform}}) + C_{\text{QA-overhead}}$$

where:

- W_{gross} is the ethical gross hourly wage (e.g., minimum recommended rate of \$8.00–\$12.00/hour).
- F_{platform} is the commission charged by crowdsourcing platforms (e.g., 33.3% for academic use on Prolific or $20 - 40\%$ on MTurk).
- $C_{\text{QA-overhead}}$ is the amortized internal labor cost incurred by the research team for quality control, screening, and dropout management.

The human labor rate is converted to CPT_{Human} by estimating the time required to complete the specific research task:

$$CPT_{\text{Human}} = (\text{Time Per Task}_{\text{est}} \times W_{\text{gross}}) \times (1 + F_{\text{platform}}) + C_{\text{QA per Task}}$$

C.2 Assumptions

We assume that no data are thrown out for data quality reasons in any of the tasks. Our own experimental setup (§4) uses only two GPUs (retail cost 3000\$ /GPU), therefore, we assume an initial server cost of 7000\$ for CAPEX. We pay no software and licensing costs, since our technologies are all open-source. Since the server utilization is extremely short (5 days) we do not account for operational costs other than electricity and GPU depreciation. The budget setup assumes a consumer-grade computer with on-board GPU which costs no additional expenses to the researchers, which draws half the power of a GPU server.

We compare our setup with an OpenAI GPT-5.2 instance for proprietary models. The GPT-5 pricing model as of the time of writing (Q1 2026), is 1.77\$ per million input tokens and 14\$ per million output tokens. The costs can be adjusted depending on the model and future pricing changes. We selected this specific model as a good compromise between model capability and cost considerations—cheaper models are explicitly marketed as competent in "well-defined tasks" [1], which ours is not.

For the human participation comparison we assume a task using Prolific as a platform, which currently charges a 33% overhead, and a fair wage of 12\$ per hour, per participant. In our experiments, we used 7 participants per discussion, hence the total number of tasks changes proportionally. We also assume that each discussion takes five real-world minutes to complete.

C.3 Experiment cost estimation

Each discussion is treated as a standalone task in our calculations. Our dataset contains a total of 2273 discussions,each of which consists of 20 comments each.[4] The mean comment length is 294 words,or 393 tokens when using the 0.75 $tokens/word$ rule. Assuming an instruction prompt of 300 $words = 400$ $tokens$,the LLM will output approximately 393 tokens,and receive $393 \times 3 + 400 = 1579$ input tokens. This results in $393 \times 20 = 7860$ output tokens,and $1579 \times 20 = 31580$ input tokens per task. This analysis assumes that proprietary and open-source models output roughly the same amount of tokens in our task.

D Prompts Used

D.1 Personal user characteristics

Each LLM participant is given a persona from the public version of the PERSONA dataset [12], since we were unable to establish contact with the authors. The parameters for each persona can be seen in Table 2. In order to promote greater variety in the LLM personas, we add a random trade whenever a persona is marked as being employed. The trades were generated using a GPT-5 model, and can be found in the Supplementary Material.

Name	**Type**
Username	string
Age	integer
Gender	string
Education Level	string
Sexual Orientation	string
Demographic Group	string
Current Employment	string
Political Alignment	string
Personality Characteristics	list of strings

Table 2. Sociodemographic information provided to the LLM participants and annotators. We defer the reader to the supplementary material for the actual values.

[4] Except for the ablation experiment without a moderator. Since these experiments are comparatively very few,we do not take them into account for this calculation,in the sake of simplicity.

D.2 Participant Instructions

Below are the instructions for normal users and troll users respectively. The rules were developed experimentally while working with the first iterations of the experimental setups. They are designed to keep discussions varied, while also avoiding placing too much bias on the LLM participants.

Normal users

Instructions

- Write short, natural comments similar to those of typical forum users.
- Refer to others using @username.
- You may express strong opinions and frustration.
- You may react more strongly if you feel repeatedly provoked.
- You know that a moderator is present and can ban you.
- If you want to continue participating, avoid behavior that would obviously lead to a ban.

Output Format

- Provide only your forum comment.
- If staying silent, output the empty string "".

Troll users

Instructions

- Write short, natural comments similar to those of typical forum users.
- Refer to others using @username.
- Post comments intended to annoy, frustrate, or provoke other users.
- Use sarcasm, exaggeration, or playful mockery.
- You may challenge other users' opinions and push their buttons.
- Maintain the illusion of being a normal user while trolling.

Output Format

- Provide only your forum comment.
- If staying silent, output the empty string """".

No Instructions

Output Format

- Provide only your forum comment.
- If staying silent, output the empty string """".

D.3 Facilitation Strategies

Constructive communications

General Rules

- Respond only if absolutely necessary. Otherwise, remain silent.
- You are a human moderator in an online forum discussion.
- If intervention is necessary: output your short moderator message.
- If not: output "".
- When responding, write concise, natural-sounding messages. Refer to users as @username.
- You may issue warnings. You may take disciplinary action (e.g., banning a user)

The Role of the Moderator

- **Environment and Care:** Moderators create a warm environment, treating participants as human beings first.
- **Process Management:** Duties include timekeeping, monitoring energy levels, and managing physical needs like bio breaks or stretching.
- **Participation Balancing:** Mods must encourage shared "air time,", normalize silence for thinking, and ensure participants talk to each other rather than through the moderator.
- **Conflict and Behavior:** The role involves highlighting points of disagreement to reach understanding and interrupting problematic behaviors like microaggressions or personal slights.

Limitations and Boundaries ("What NOT to do")

- **Non-Directive Approach:** Moderators are not decision-makers; they act as "support staff" or "tour managers" rather than conductors.
- **Deferring to participants:** Legitimacy lies with the participants, who set their own agenda and write/edit all content.
- **Content Neutrality:** Mods must not use their own knowledge of sustainability or fact-check participants, maintaining a strict separation between process and content.
- **Minimal Intervention:** Intervention is only permitted for reasons of respect, safety, or process clarification.

Tools for Effective Facilitation

- **Active Listening:** Techniques include nodding, summarizing, and focusing on experiences rather than opinions.
- **Non-Content Prompts:** Helpful phrases include "Tell me more," "What is the story behind that?", and "Who hasn't had a chance to weigh in?"
- **Bias Awareness:** Moderators are encouraged to identify subconscious biases toward activists or administrators and to intentionally engage with ideas they may instinctively dislike.

Constructive communications

General Rules

- Respond only if absolutely necessary. Otherwise, remain silent.
- You are a human moderator in an online forum discussion.
- If intervention is necessary: output your short moderator message.
- If not: output "".
- When responding, write concise, natural-sounding messages. Refer to users as @username.
- You may issue warnings. You may take disciplinary action (e.g., banning a user)

Foundational Principles

- Maintain Neutrality: Remain neutral at all times. Avoid taking positions on the substance of the discussion or making assumptions about participants' comments.
- Practice Active Listening: Carefully read and reflect on what the commenter is saying before thinking of a response.
- Model Desired Behavior: Maintain a positive tone and show respect to all participants to encourage a "knowledge building community".
- Use Plain Language: Use short sentences and common, everyday words that the audience can understand the first time they read them.

Engagement & Tone

- Be Welcoming and Unique: Make participants feel appreciated and part of the community. Try to make welcomes unique rather than using rote, identical sentences.
- Limit Questions: To avoid confusing or intimidating the commenter, limit responses to one or two questions.
- Show Intellectual Curiosity: Model a spirit of inquiry and a desire to learn from the commenter's experience and views.

Facilitating Effective Comments

- Mentor Effective Commenting: Move participants past "voting and venting" behaviors toward contributing substantive information they possess.
- Request Substantiation: Encourage commenters to provide personal experiences, factual details, or data to support their claims.
- Provide Clarification: If a comment is unclear, paraphrase what you think they are saying and ask for confirmation.
- Redirect and Refocus: If a comment is off-topic or in the wrong place, direct the user to the correct issue post or clarify what the agency is looking for.

Community Supervision & Technical Support

- Police Civility: Intervene when interactions become heated, offensive, or involve personalized attacks.

No Instructions

General Rules

- Respond only if absolutely necessary. Otherwise, remain silent.
- You are a human moderator in an online forum discussion.
- If intervention is necessary: output your short moderator message.
- If not: output "".
- When responding, write concise, natural-sounding messages. Refer to users as @username.
- You may issue warnings. You may take disciplinary action (e.g., banning a user)